

Deep Learning Theory and Applications

# Generative Adversarial Networks (GANs)

Yale

CPSC/AMTH 663

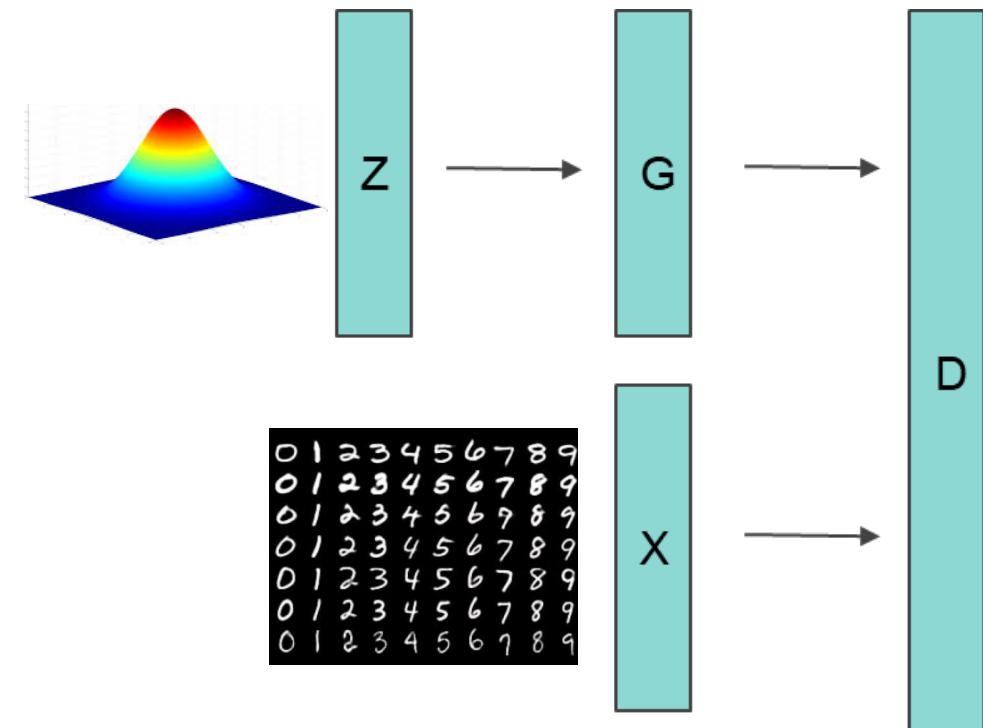




# Outline

1. GANs
2. Conditional GANs
3. Autoencoder GANs
4. DiscoGANs
5. Cycle GANs
6. MAGAN
7. TraVeLGAN

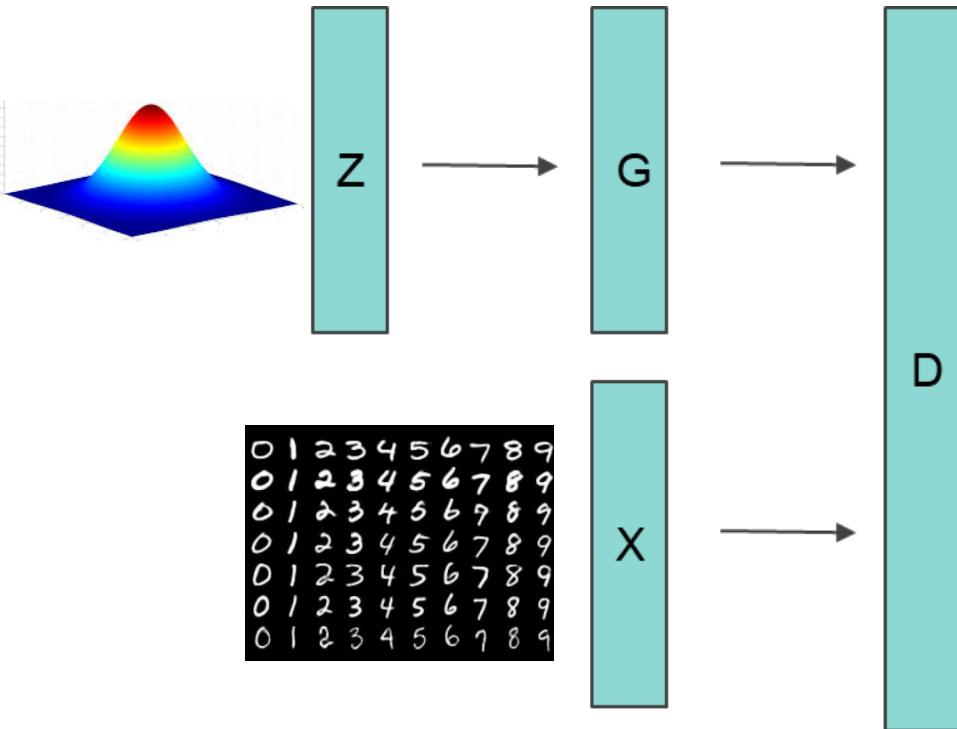
# GANs



- Sample from  $Z$ 
    - Easy to do!
  - One network learns a mapping  $G(Z)$
  - Second network tries to distinguish true examples of  $X$  from “fake” samples  $G(Z)$



# GANs

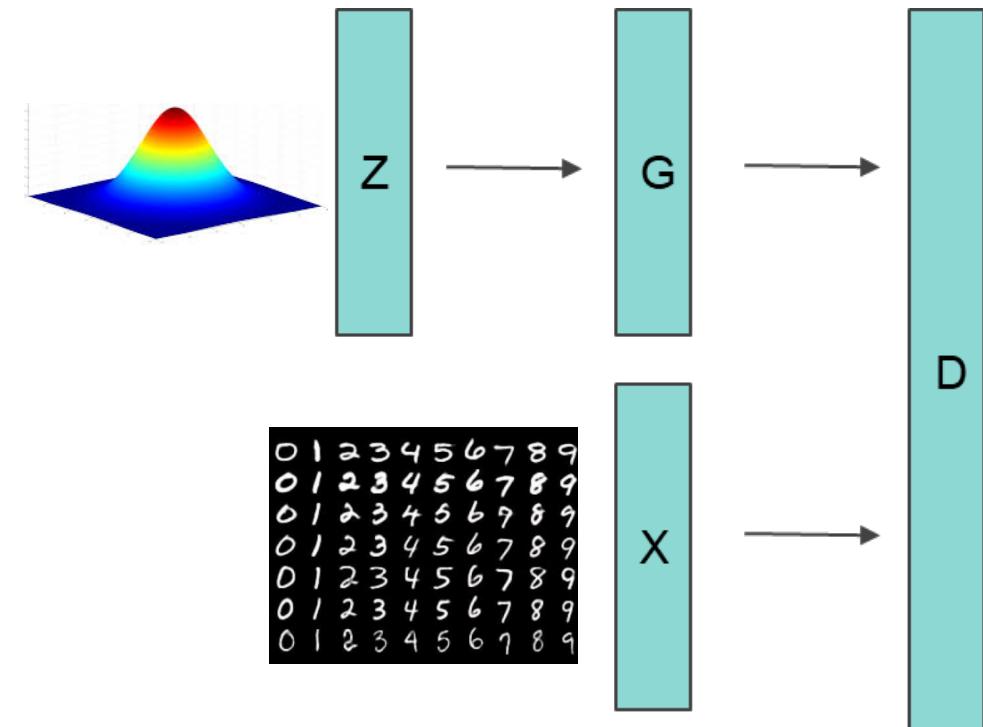


- After training, we throw out  $D$
- Sample from  $Z$
- Compute  $G(Z)$
- New samples from an estimate of  $P(X)$  (not  $X$ )!
- Note that the GAN did not just memorize the finite data set, as the generated samples do not appear exactly in  $X$

$G(Z)$	<table border="1"><tr><td>9</td><td>9</td></tr><tr><td>0</td><td>0</td></tr><tr><td>2</td><td>2</td></tr><tr><td>8</td><td>8</td></tr></table>	9	9	0	0	2	2	8	8	Closest $x$ to $G(Z)$
9	9									
0	0									
2	2									
8	8									



# GANs



- Training alternates between D and G
- Discriminator tries to classify true samples as 1 and fake samples as 0
- Generator tries to fool the discriminator
- Gradients flow in between the networks



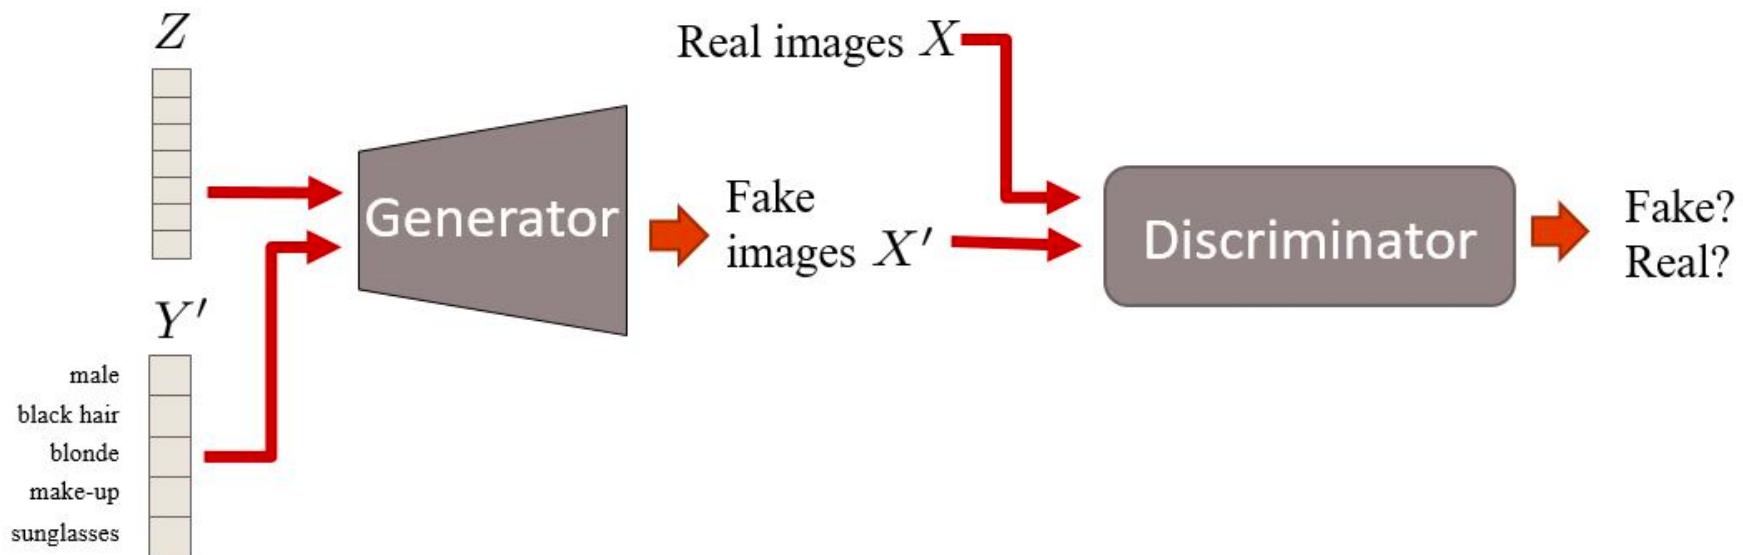
# Conditional GANs

- GANs take noise  $z$  and turn it into our data  $x$
- Great! But what if we want to control it?
  - “It generated *another* image of a 3? I wanted a 7 this time!”
- Conditional GANs:
  - If we have labels, we want to teach it to generate not just any  $x$ , but an  $x$  of a particular label when given some noise  $z$



# Conditional GANs

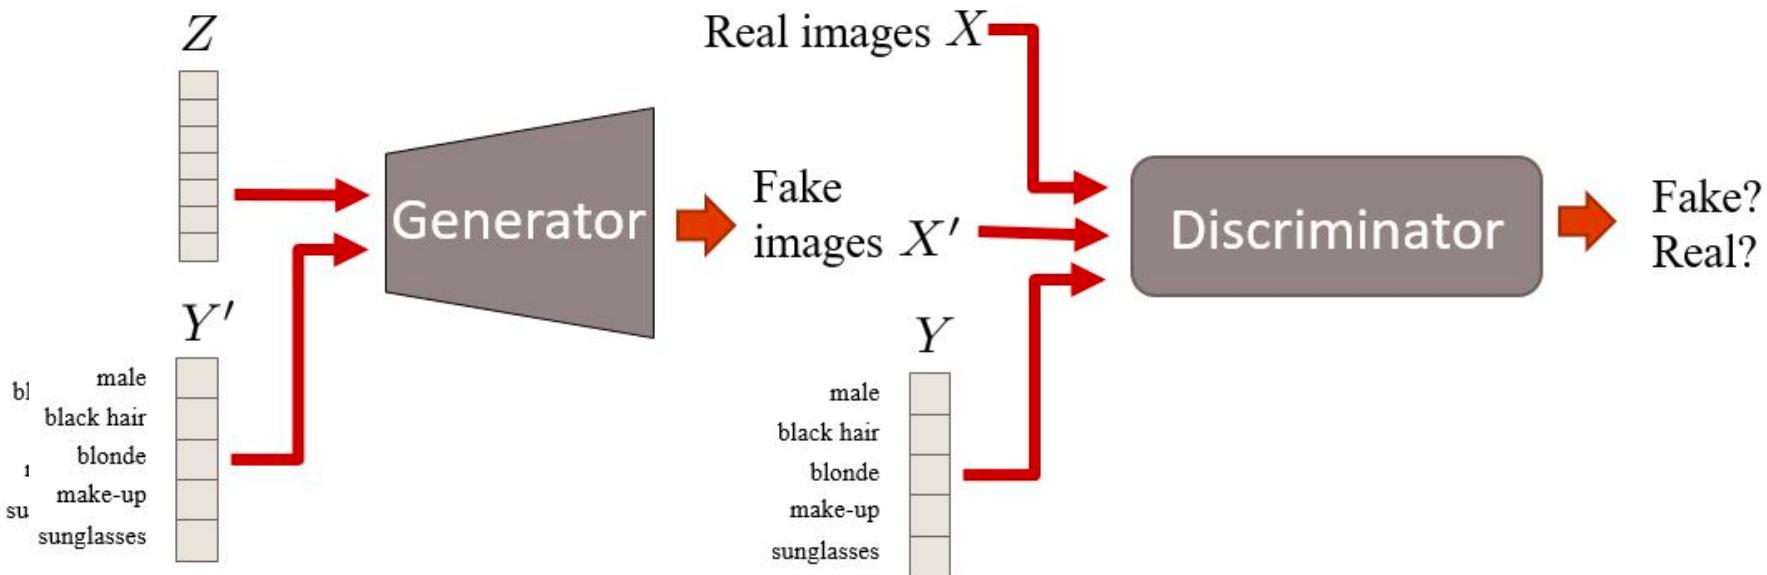
- Feed the generator both a sample from  $z$  and the condition
- Discriminator tries to tell real/fake apart as before
  - The generator could generate really realistic images that don't match the condition and the discriminator would never know





# Conditional GANs

- Have to also give the discriminator the labels for each point
  - The discriminator not only looks to see if it is realistic, but also if it is an image/condition pair that doesn't appear in real data



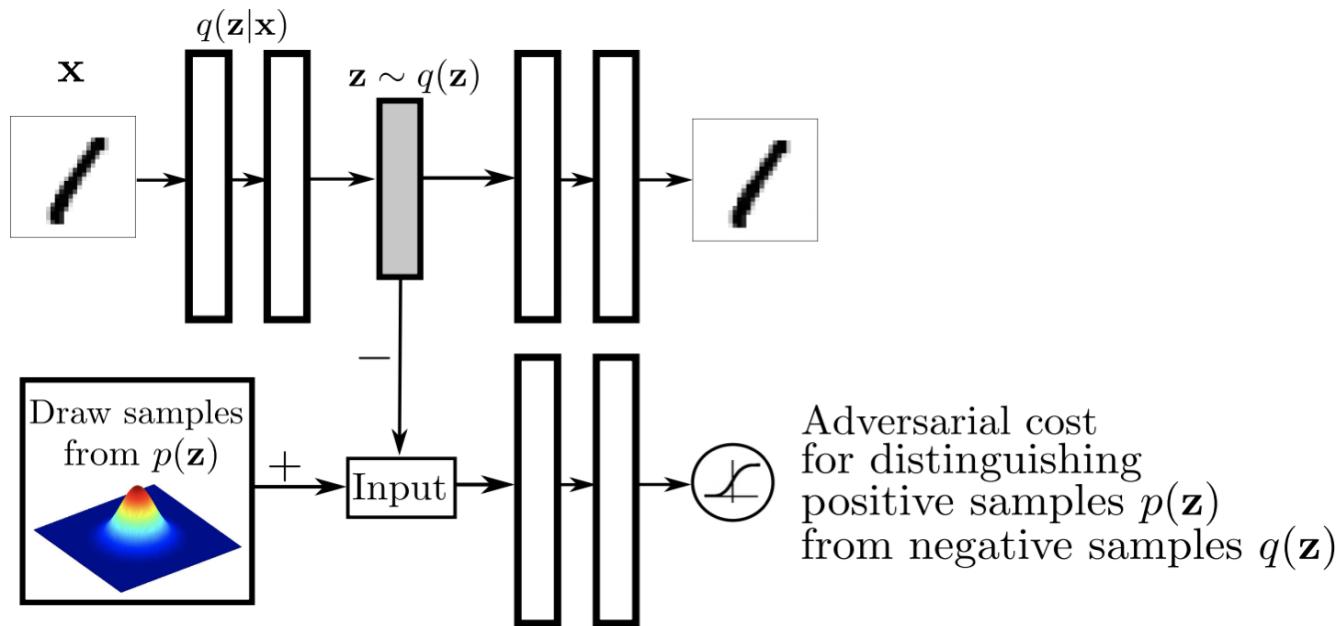


# Can adversarial training be used on Autoencoders ?

- Can we use it to make an improved version of a VAE?
- Can the role of the KL-divergence penalty be replaced by an adversary?



# Adversarial Autoencoders

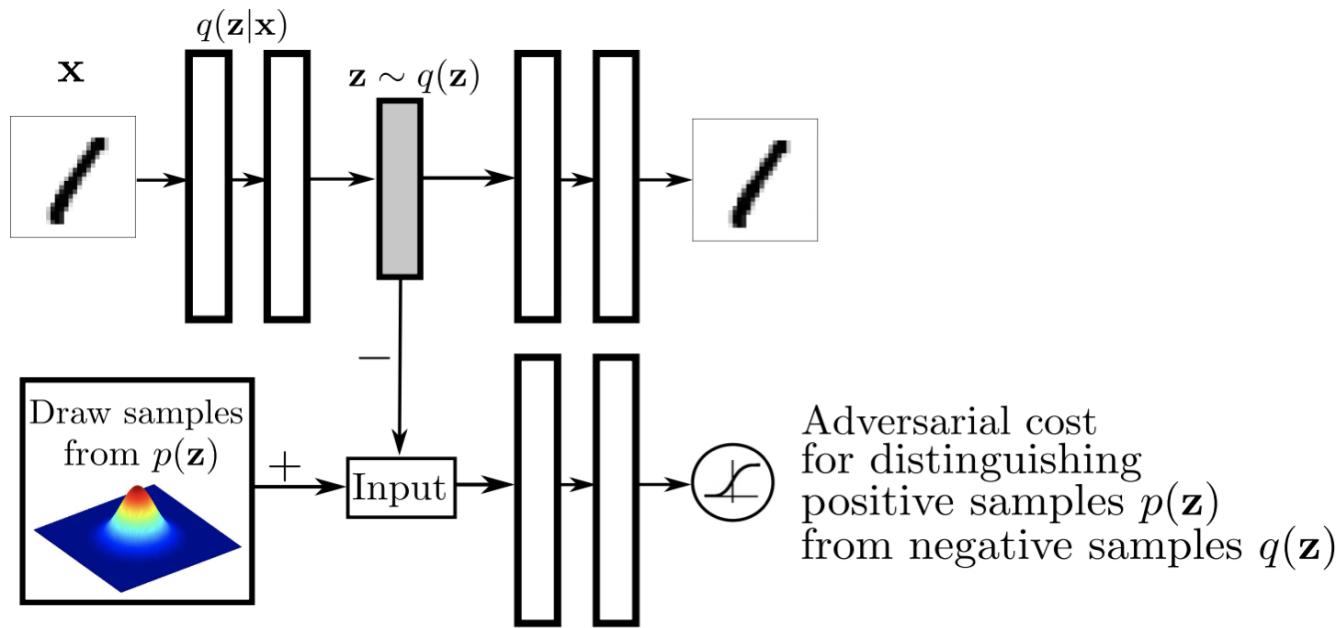


## Differences

1. Discriminator classifies samples from  $\mathbf{z}$  against the AE hidden layer
  - a) Low-dimensional and easy distributions
2. Generator has both encoder and decoder
3. Generator receives gradient signal from discriminator **and** reconstruction loss



# Adversarial Autoencoders



- Generate new sample as follows:
  1. Sample from  $p(\mathbf{z})$
  2. Decode  $\mathbf{z}$
- Encoder and discriminator only needed for training, not sampling



# Other advantages of AANs

- The discriminator has a simpler job: just needs to distinguish between latent layer distribution and gaussian distribution
- Has denoising and other advantages of an autoencoder



# DiscoGAN: Motivation

- GANs and adversarial AEs sample from one domain
- What if we have two domains, A and B?
- Can we learn to turn a point in A into a point that looks like its from B?
- Other attempts required paired input  $(a_i, b_i)$ 
  - Uses are limited since labels can be expensive or difficult to obtain
- Can we do this with an unsupervised training setup?



# DiscoGAN: Motivation

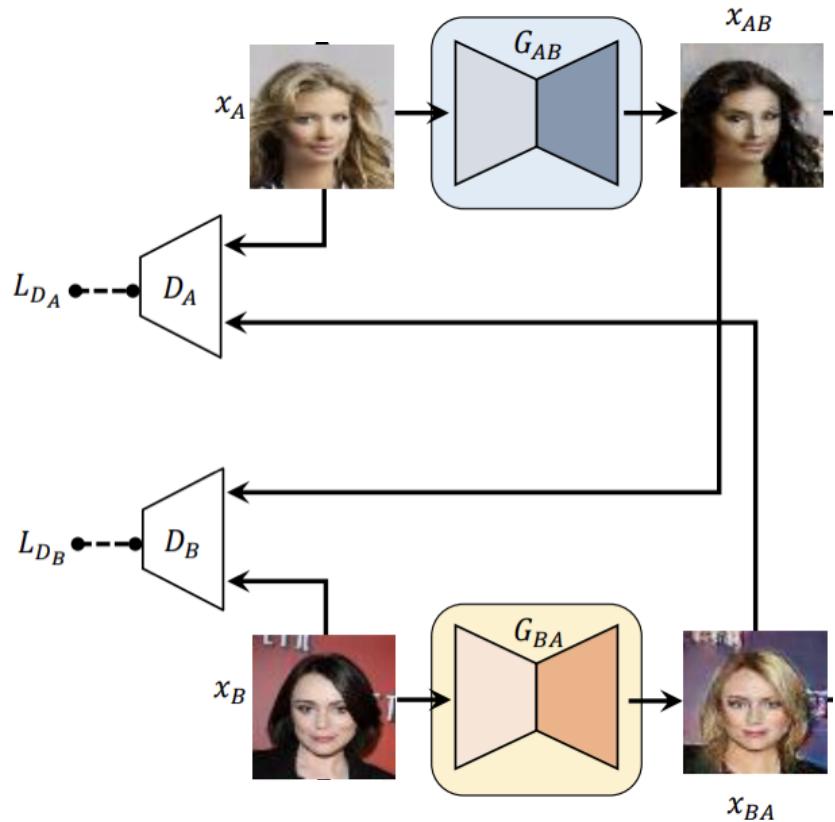
- Domain A: blond hair



- Domain B: black hair



Previous approach



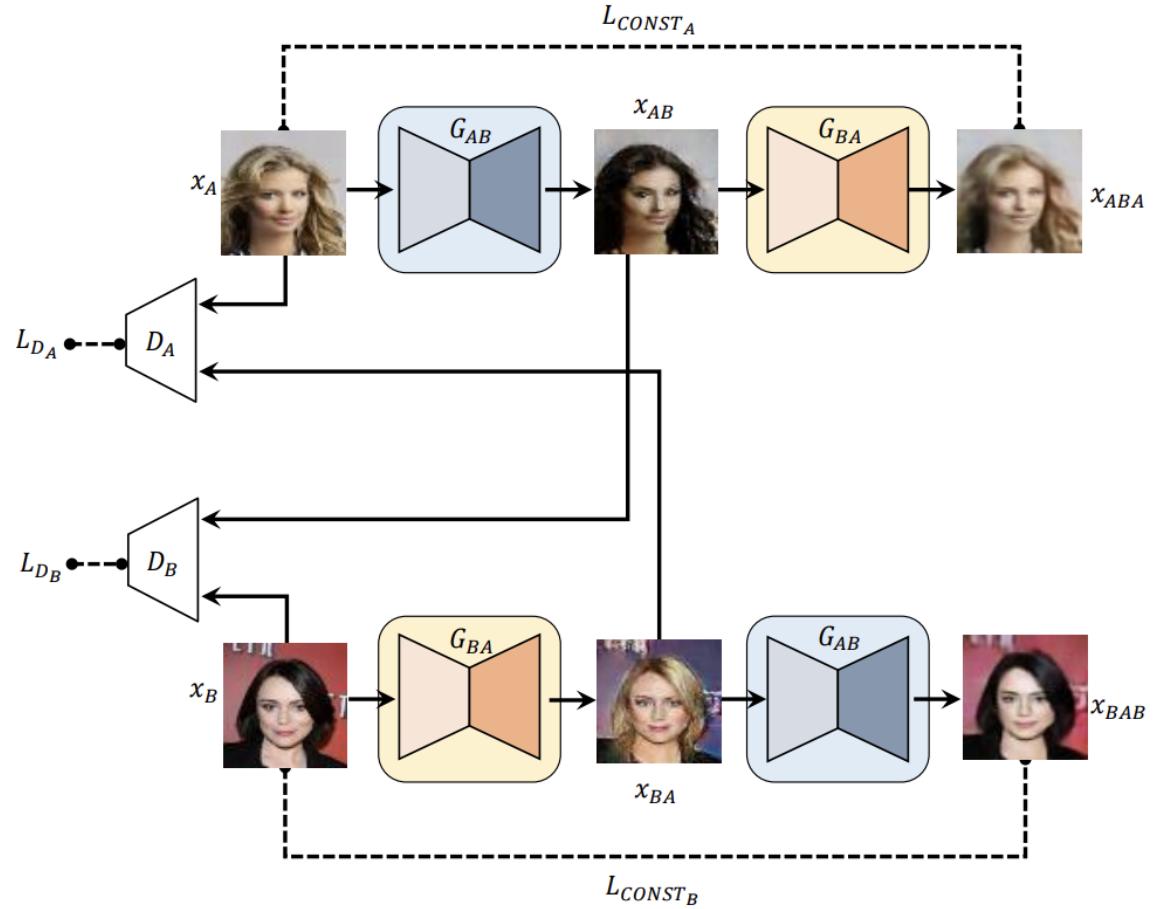


# DiscoGAN

- Domain A: blond hair



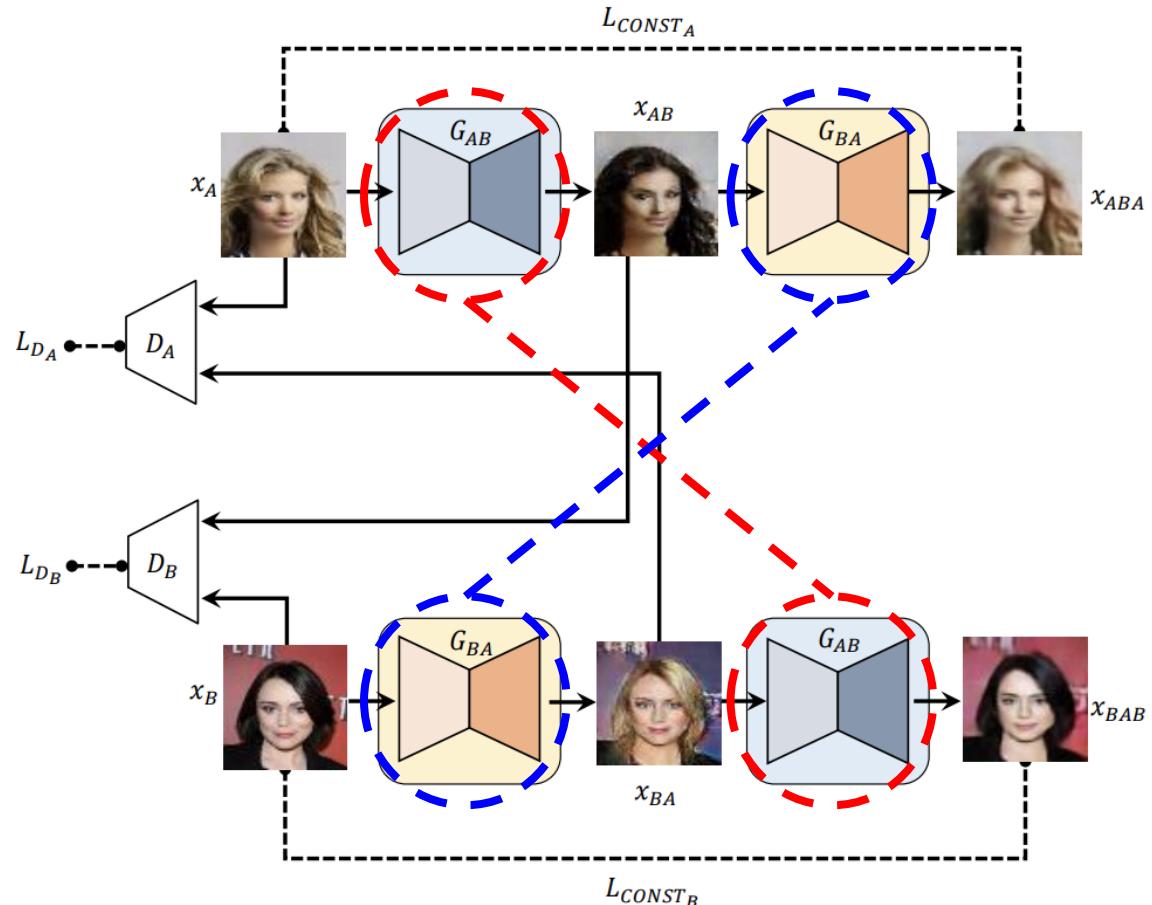
- Domain B: black hair





# DiscoGAN

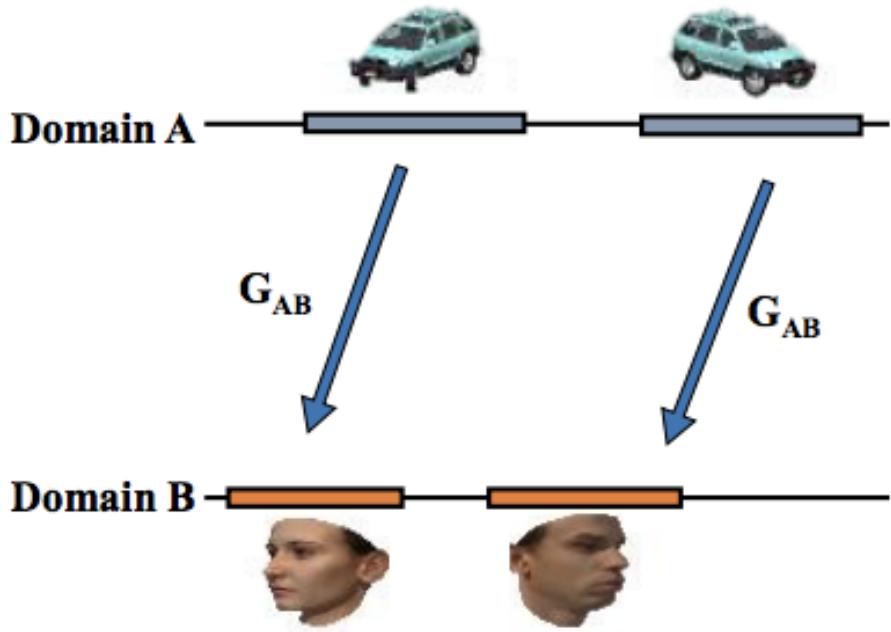
- Weight sharing/reuse is important
  - Same reason CNNs work
- Same network weights/activations must perform multiple tasks
  - Indirectly increases number of training samples and generalizability
- Reduces number of weights to train
  - Optimization is easier





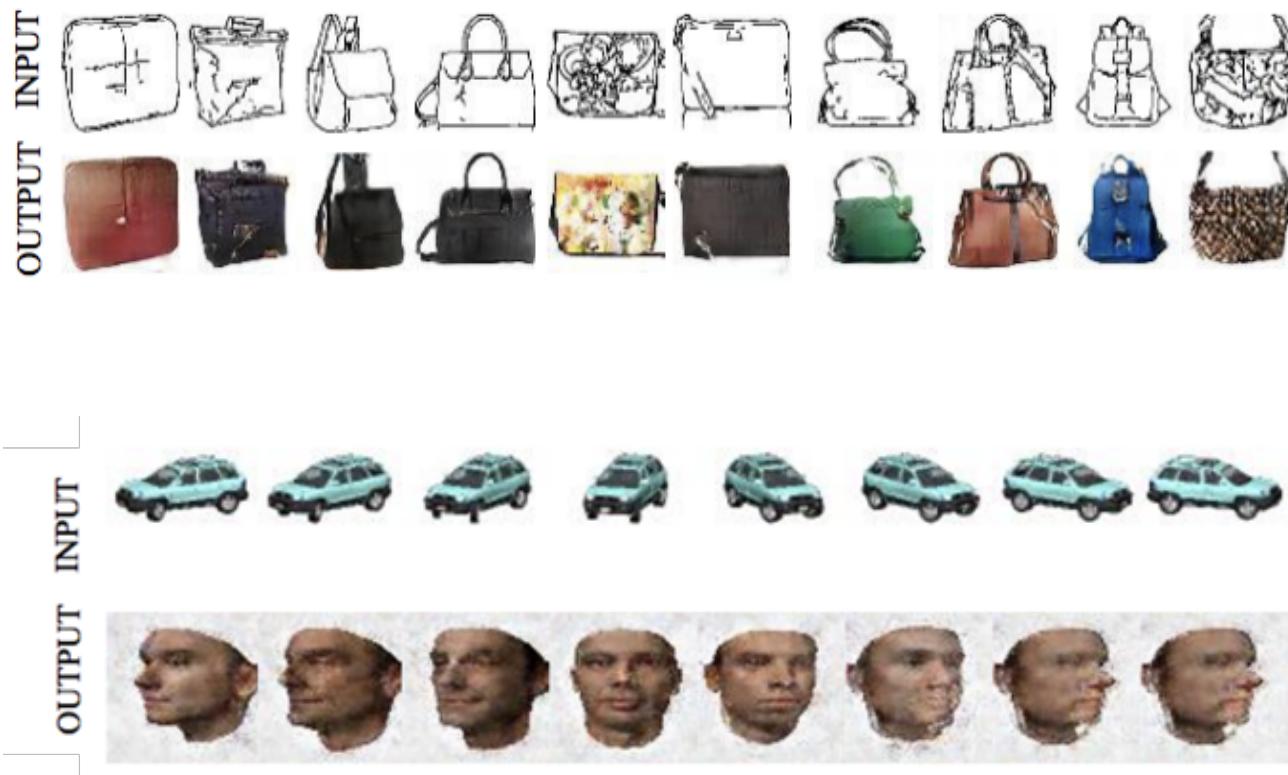
# DiscoGAN

- Weight sharing/reuse is important
  - Same reason CNNs work
- Same network weights/activations must perform multiple tasks
  - Indirectly increases number of training samples and generalizability
- Reduces number of weights to train
  - Optimization is easier





# DiscoGAN





# DiscoGAN

## Issues

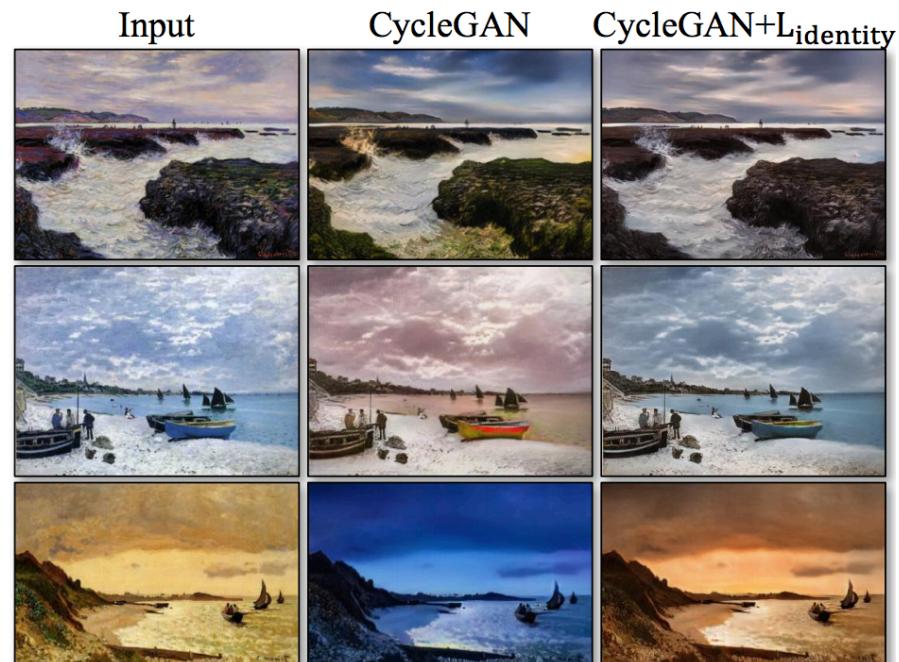
- Is this the correct shoe for this purse?
- Is there a better match?
- Original paper doesn't talk about it
- Sometimes the mapping isn't correct





# CycleGAN

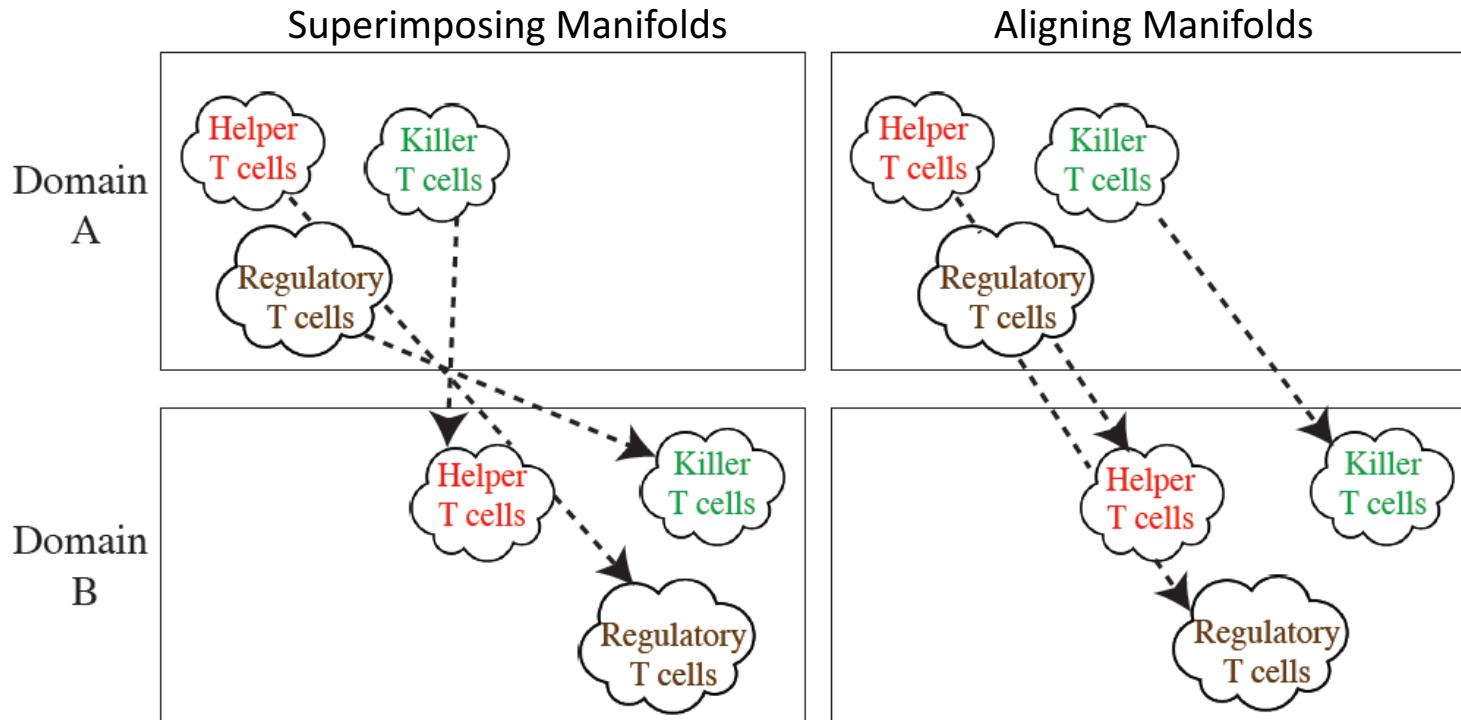
- Came out shortly after DiscoGAN, popularized the idea
- Added another component to the loss:
  - Normally  $G_{AB}$  takes  $x_A$  to domain B
  - This means  $G_{AB}(x_A)$  will be different from  $x_A$
  - Add additional constraint that  $G_{AB}(x_B) = x_B$ 
    - Intuitively, G should take something not in A to B, but if it's already in B, it shouldn't change it





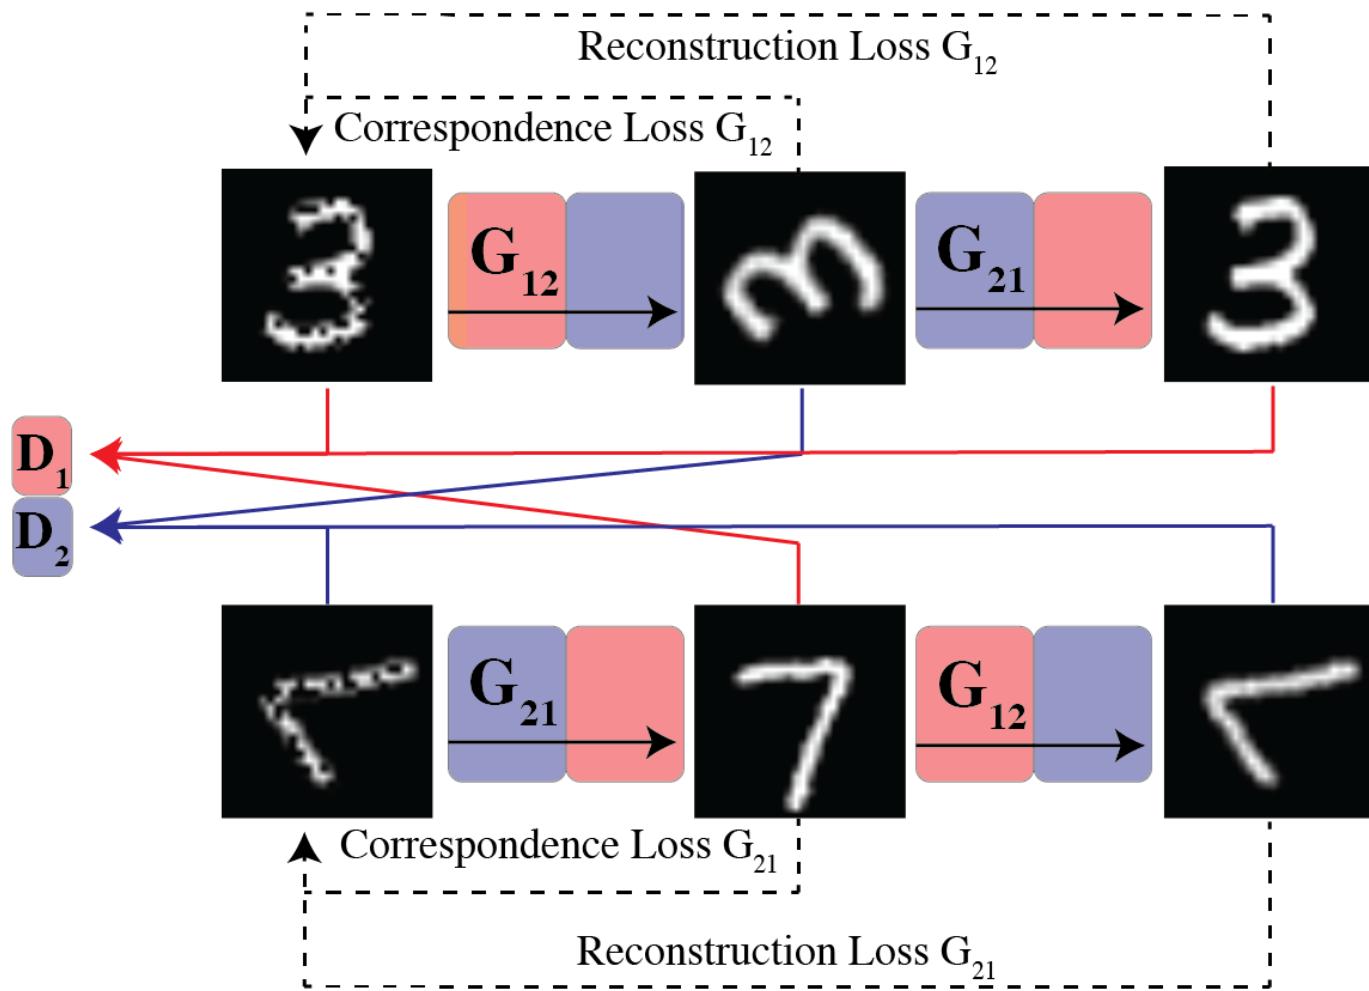
# Manifold Alignment GAN (MAGAN)

- Amodio & Krishnaswamy (2018)
- Consider manifolds A and B to be two CyTOF panels
  - Different markers, samples from same tissue
- Panel A measures 40 markers
- Panel B measures 40 markers
- 10 shared markers and 30 unique markers, each
- Using 10 shared markers, find corresponding cells across panels





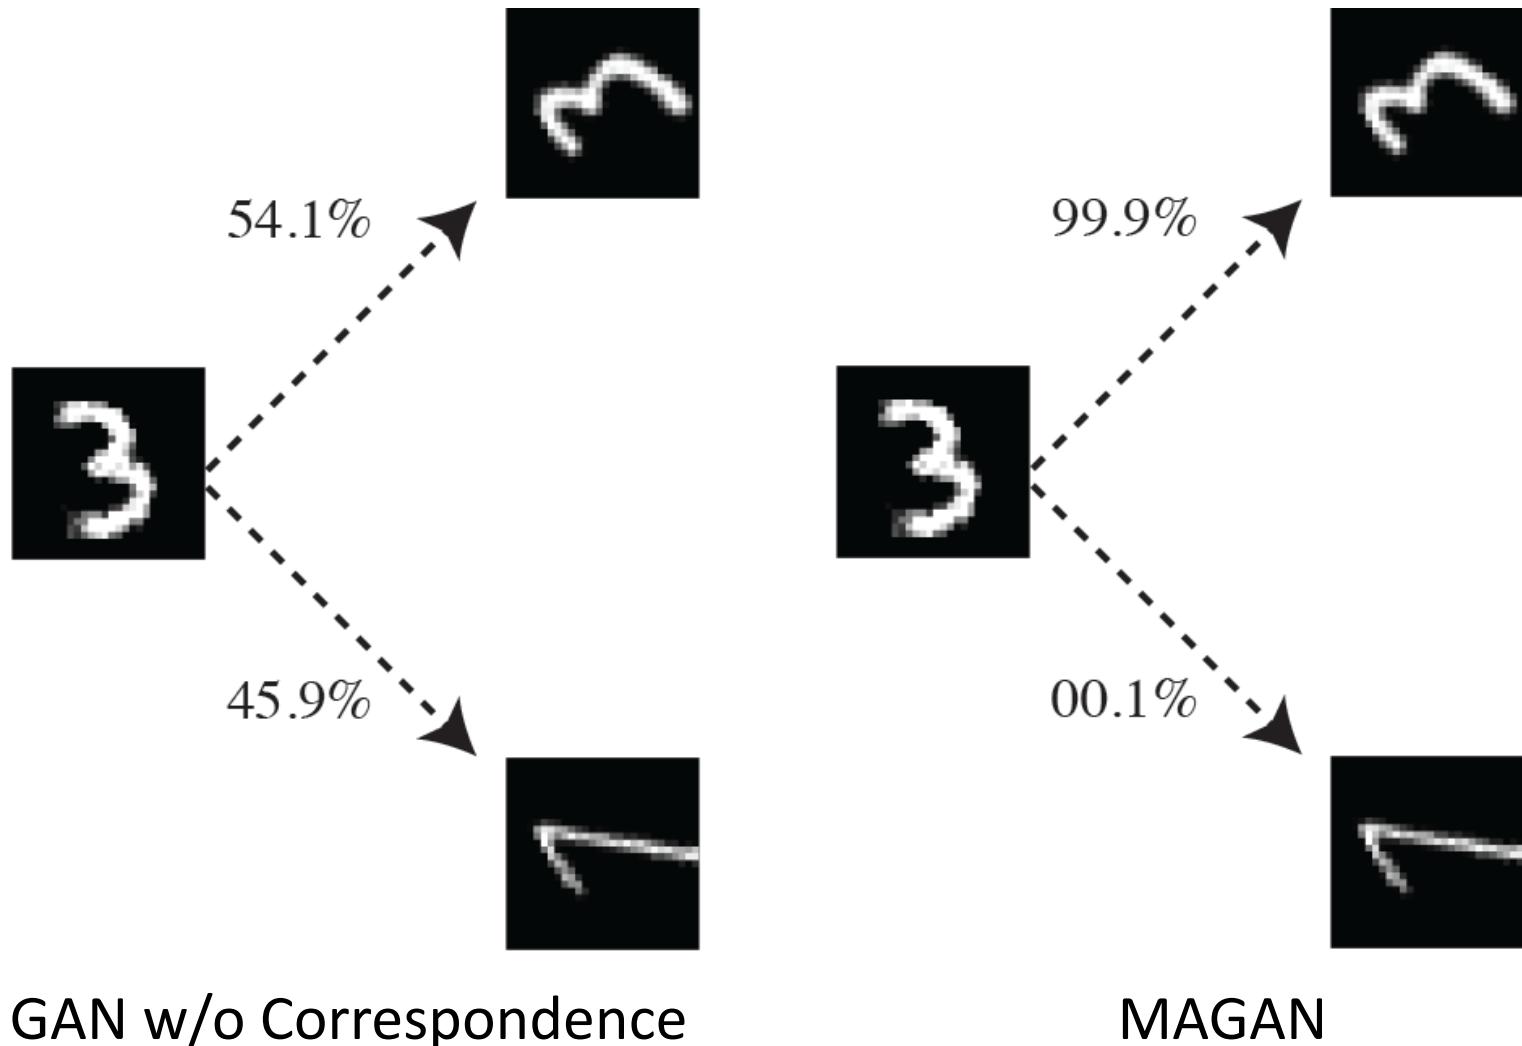
# MAGAN Architecture





# Matching rotated digits

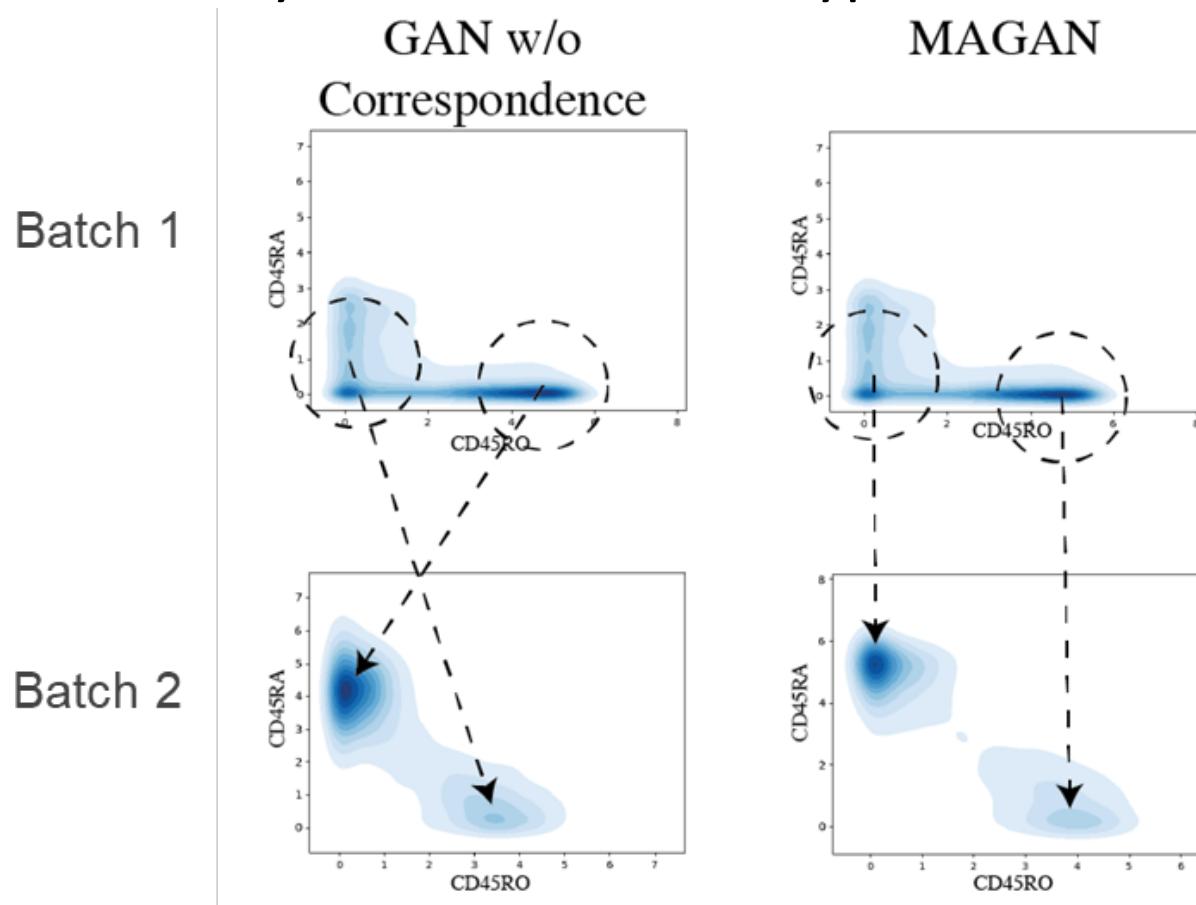
- 100 simulations of each model





# Aligning CyTOF Batches w/ MAGAN

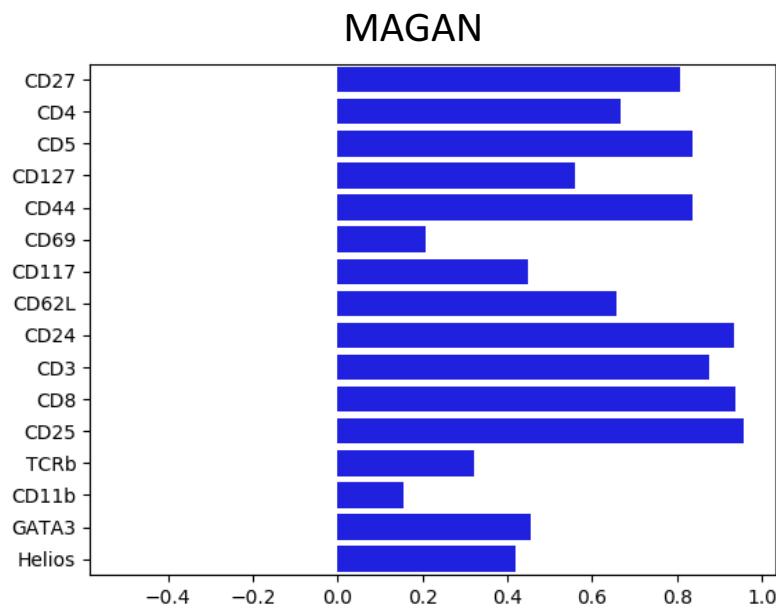
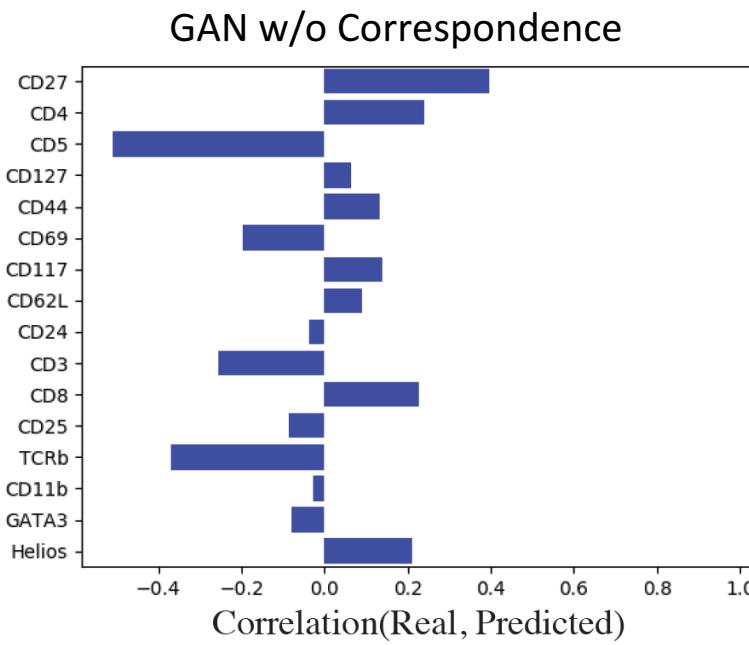
- Two underlying cell populations measured in each batch (naïve T cells and central memory T cells)
  - Dropout in CD45RA in batch 1
- MAGAN correctly matches the cell types in the two batches





# Aligning CyTOF Panels w/ MAGAN

- Validation on publicly available CyTOF data from developing mouse brain\*
- Hold out a shared marker in one panel
- Apply MAGAN to reconstruct held-out marker
  - Compare correlation between real value and prediction



\* Setty, Manu et al. "Wishbone identifies bifurcating developmental trajectories from single-cell data". *Nature Biotechnology*.

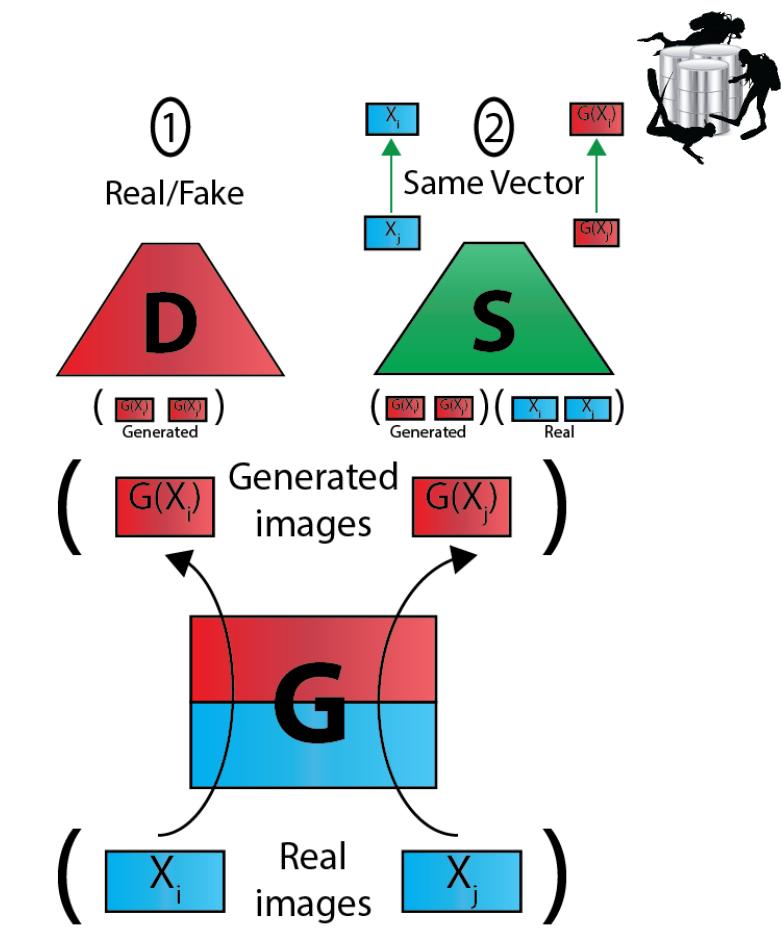


# Limitations of Cycle-consistency

- Can only learn simple **style-transfer** transformations
  - e.g. paint a horse to look like a zebra
- Can only learn **one** transformation
  - e.g. all horses become zebras, and nothing else is transformed
  - Thus domains must be homogenous
  - Papers do a lot of cropping and centering...
- Transform based on the **pixel space**
  - Is MSE between two 256x256x3-dimensional images meaningful?
- Assumes **any invertible** transformation is meaningful
  - What does invertibility have to do with semantic meaning?
- Not **interpretable**
  - Why did individual  $i$  in domain 1 get paired with this particular individual in domain 2? CycleGAN offers no insight into its decision

# TraVeLGAN

- Instead of enforcing cycle-consistency, enforce that the vector between two real points is the same as the vector between their corresponding two generated points
  - Key: do this in a latent space, not pixel space!
  - Where do we get this latent space?
- Learn another network (**S**) to go along with G and D to find a latent space where these transformations are preserved
  - Train S alongside the other networks, gradually developing an appropriate semantics
- Since we are no longer enforcing pixel-wise MSE with cycle-consistency, we have a lot more freedom to change the image for domain transfer now



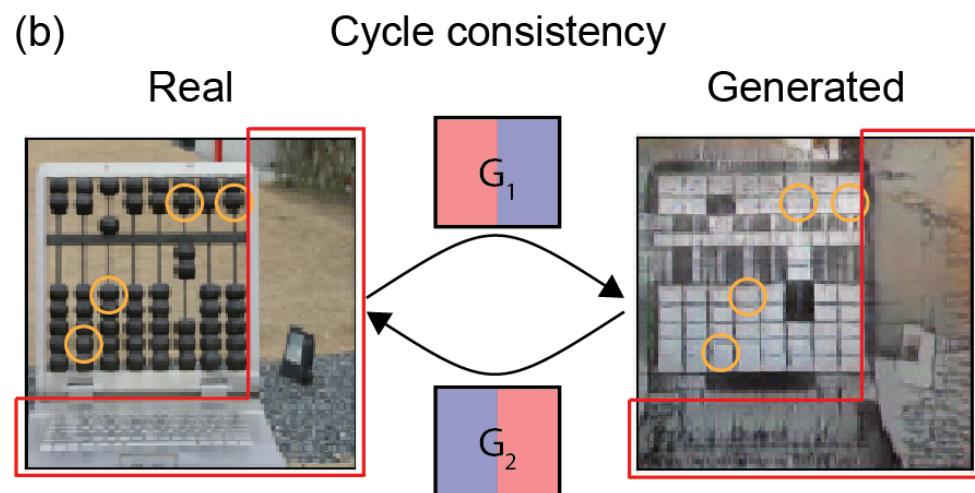
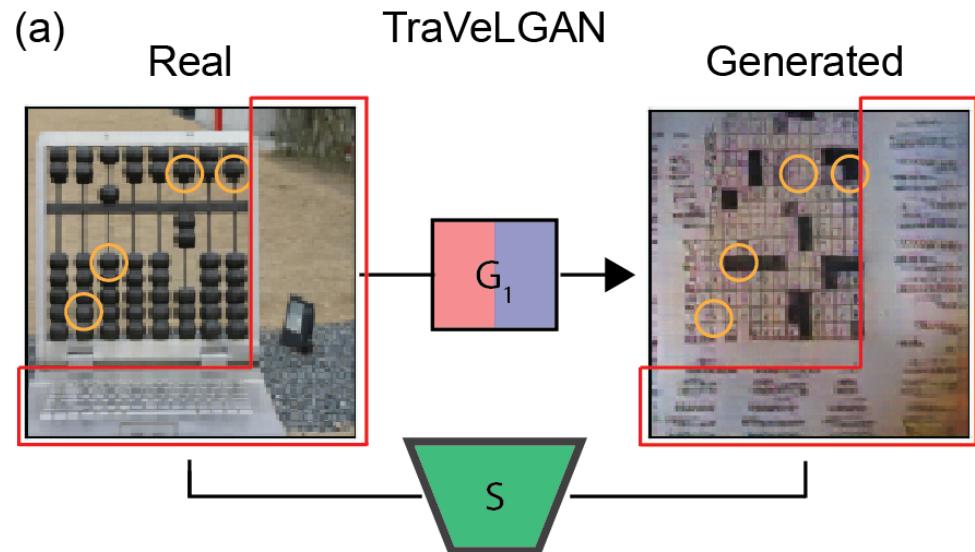
**Transformation Vector Learning  
(TraVeLing)**

$$\begin{array}{c}
 \boxed{G(x_i)} \quad ? + \left( \boxed{x_i} - \boxed{x_i} \right) = \boxed{?} \\
 \downarrow \\
 \boxed{G(x_i)} \quad + \left( \boxed{x_i} - \boxed{x_i} \right) = \boxed{G(x_i)}
 \end{array}$$



# TraVeLGAN: examples

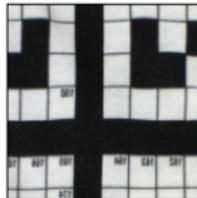
- Transferring abaci to crosswords
- The desired function is not invertible
- Any crossword configuration is a valid abacus configuration, but the reverse is not true
- Changes to the background of an image usually have to be lossy, which we can't have if we require pixel-wise cycle-consistency



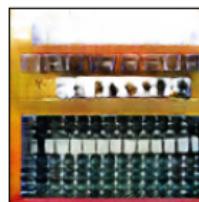


# TraVeLGAN: examples

Original



TraVeL  
GAN



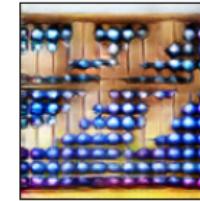
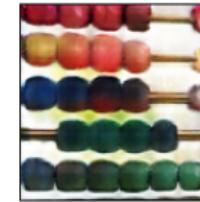
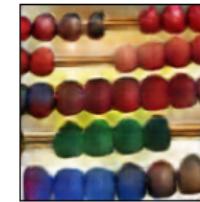
Cycle



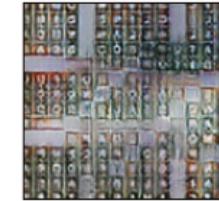
Original



TraVeL  
GAN

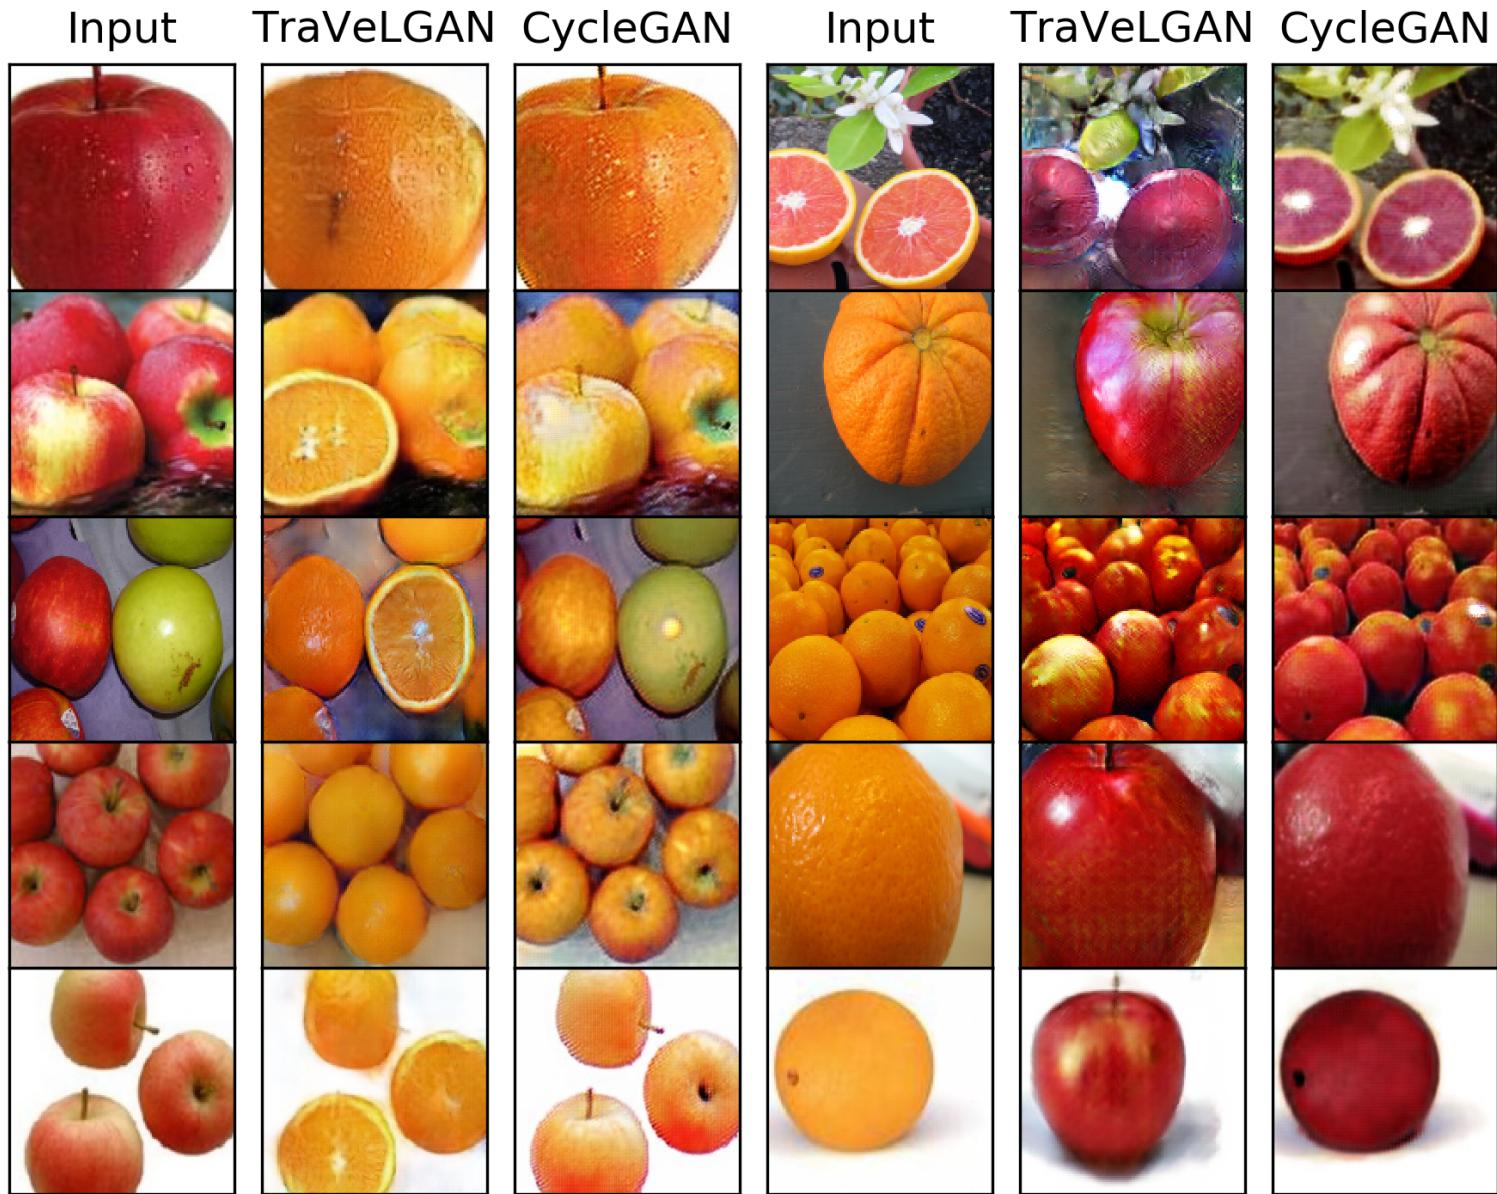


Cycle





# TraVeLGAN: examples





# Further reading

- Goodfellow et al., Section 20.10.4
- Lectures on GANs CS 11-785 at CMU
- DiscoGAN: <https://arxiv.org/abs/1703.05192>
- CycleGAN: <https://arxiv.org/abs/1703.10593>
- MAGAN: <http://proceedings.mlr.press/v80/amodio18a.html>
- TraVeLGAN: <https://arxiv.org/abs/1902.09631>