

Deep Learning Theory and Applications

Universality of Neural Networks

Yale



Outline





1. Introduction
2. Universality with one input and one output
3. Many inputs
4. Wrap-up
 - Beyond sigmoid functions
 - Fixing step functions

Universality of neural networks




- Neural networks can compute (almost) any function
- No matter the function, there is guaranteed to be a neural network that for every possible input x , the network closely approximates $f(x)$



Universality of neural networks



- Also applies to functions with many inputs and many outputs



Universality of neural networks



- In other words, neural networks have a ***universality***
 - No matter what function we want to compute, there is a neural network that can do it
- This is true for networks with only a single hidden layer
- Most proofs are quite technical
- Today, we'll give a simple and visual explanation of the universality theorem

A note on universality





- Universality theorems are quite astonishing
 - The ability to compute an arbitrary function is remarkable
- Almost any process can be thought of as function computation
- Examples:
 - Name a piece of music based on a short sample
 - Translate Chinese text to English
 - There may be many possible functions
 - Generate a plot description from a movie file
- Universality means that neural networks can do all of these things and more
 - However, this doesn't mean that we have good techniques for constructing or even recognizing such a network


Some caveats



- Neural networks can't *exactly* compute any function
 - Rather we can get an *approximation* that is as good as we want
 - Increasing the number of hidden neurons can improve the approximation



Poor approximation



Better approximation

- Can do better with even more hidden neurons

Some caveats



- Let's make this more precise
- Suppose we're trying to approximate $f(x)$ within some accuracy $\epsilon > 0$
- The guarantee is that with enough hidden neurons, there exists a neural network whose output $g(x)$ satisfies $\forall x$

$$|g(x) - f(x)| < \epsilon$$

Some caveats




- Second, we can only guarantee this accuracy for *continuous* functions
 - If a function is discontinuous, then it won't be generally possible to approximate it at each point since the neural network output is continuous
- However, often a continuous approximation of a discontinuous function is good enough

Some caveats



- So in practice, continuity isn't a major limitation
- Summary: neural networks with a single hidden layer can be used to approximate any continuous function to any desired precision
 - For simplicity, we'll focus on the case with 2 hidden layers


Universality with one input and one output



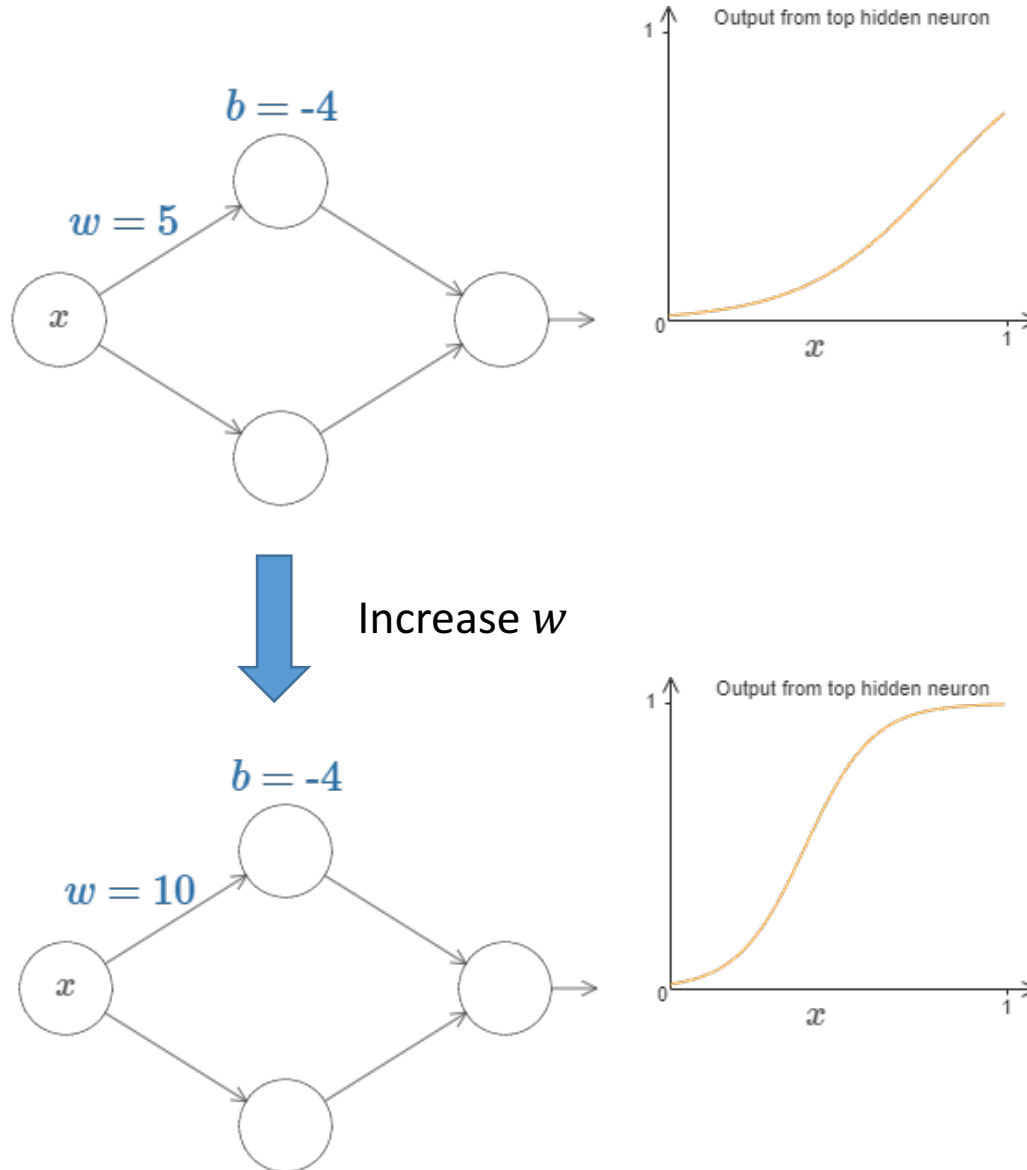
One input, one output




- Start simple



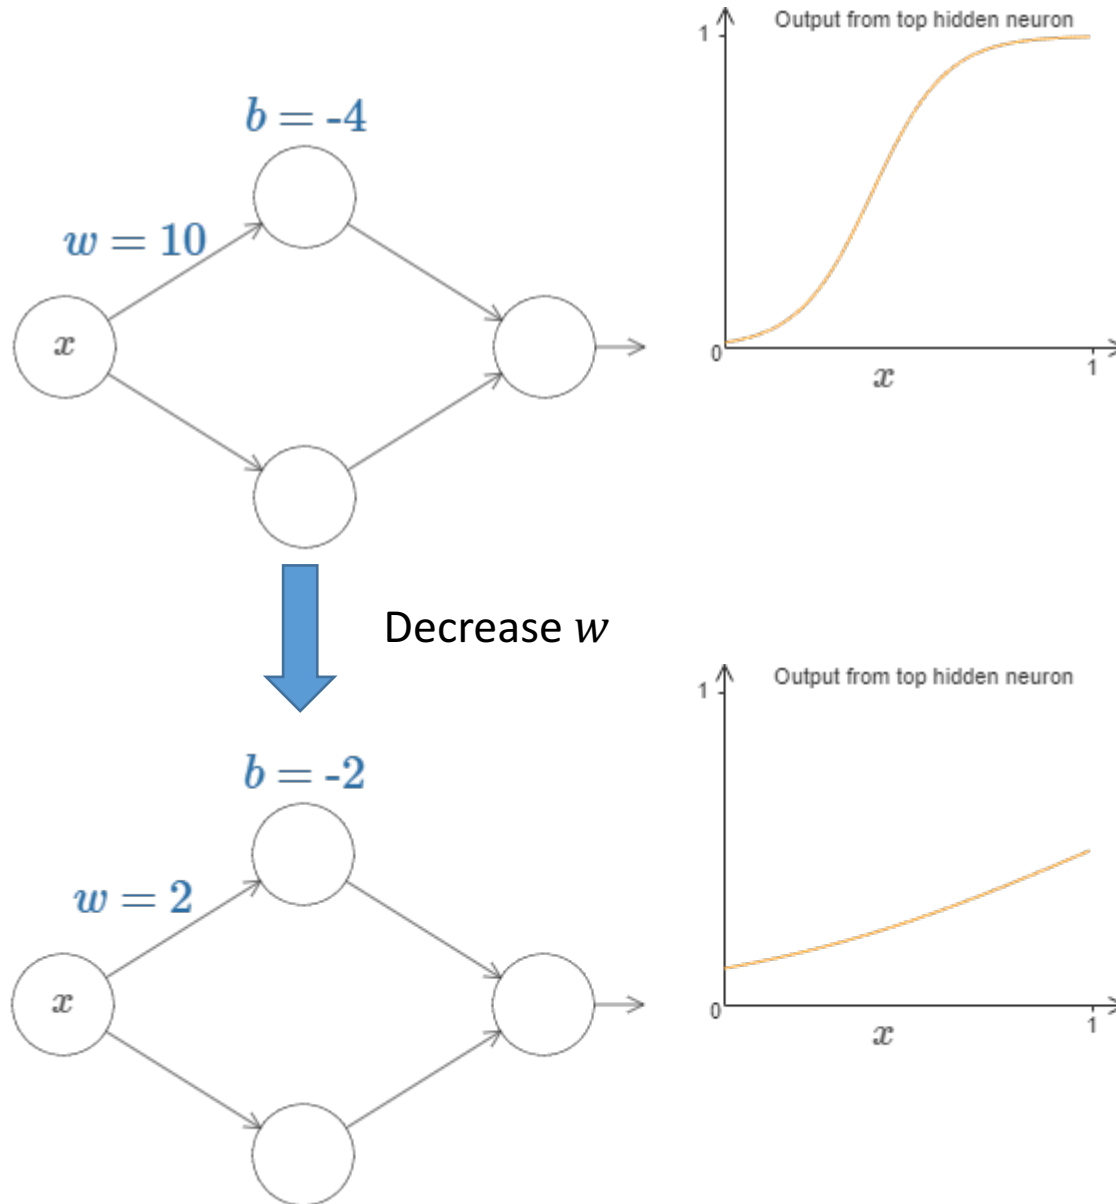
One input, one output



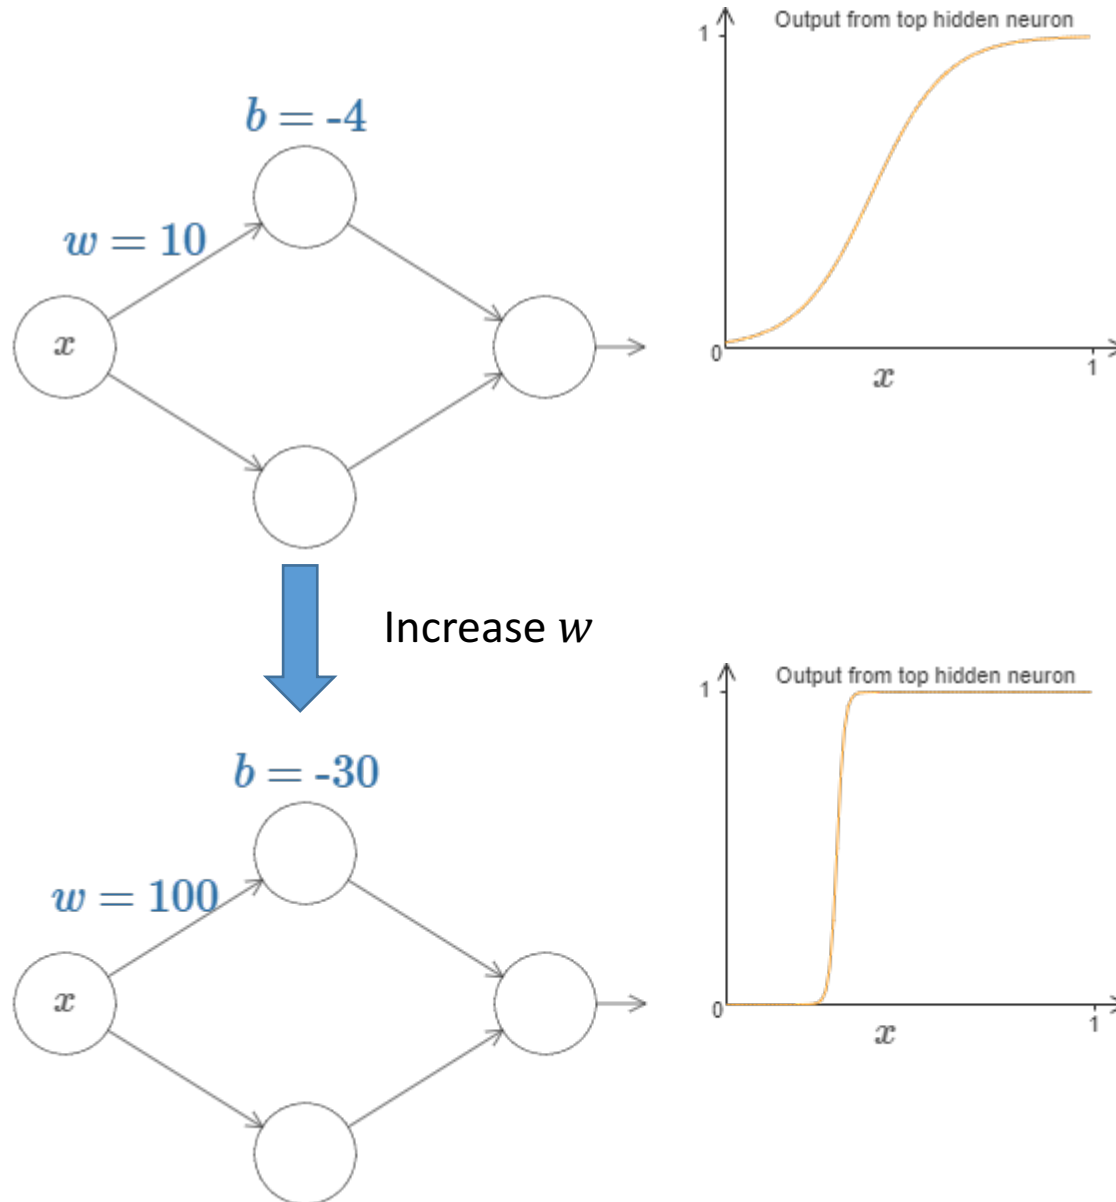
One input, one output



One input, one output





One input, one output



One input, one output





- We can simplify our analysis a lot by using a step function
 - The output layer is a sum of contributions from all hidden neurons
 - Easier to analyze the sum of step functions
 - Approximate a step function by setting w to be very large, and modifying the bias appropriately
 - Later, we'll cover the effect of this approximation



One input, one output



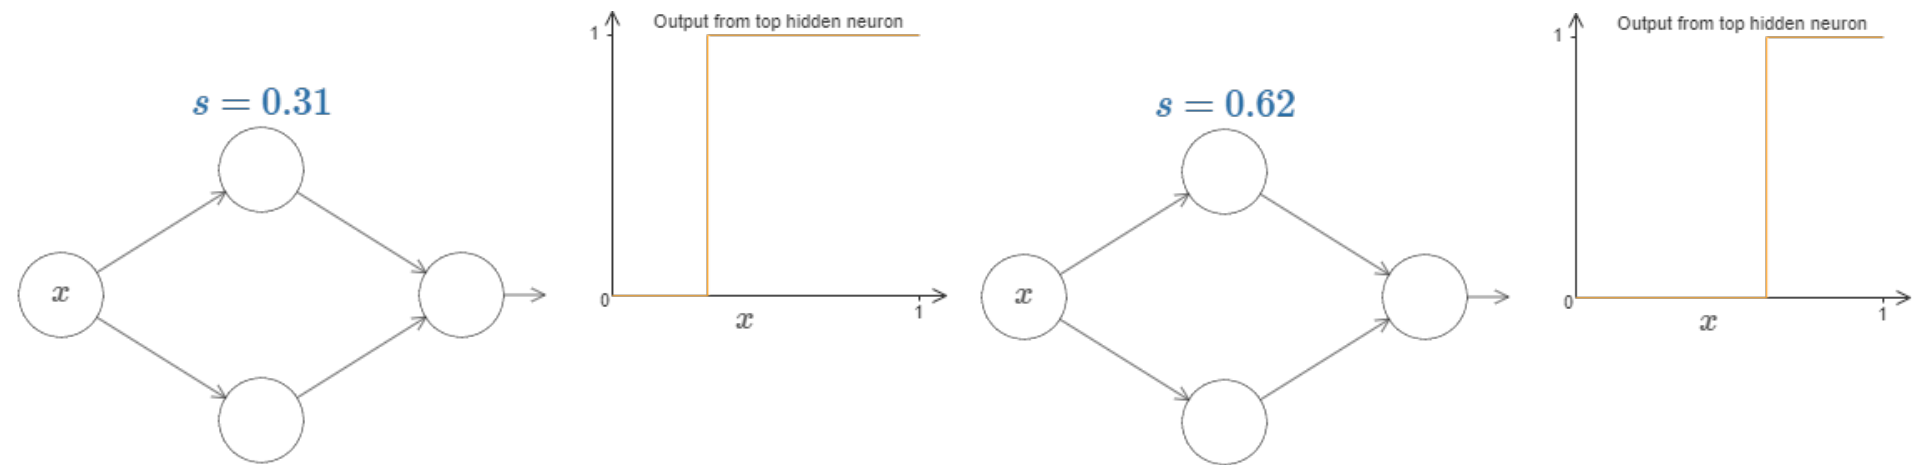
- Where does the step occur?
- The position of the step is proportional to b and inversely proportional to w
 - The step is at position $s = -\frac{b}{w}$



One input, one output



- We can simplify things by using a step function with parameter s
 - I.e., we set w to be some very large value and then adjust b
 - Recover $b = -ws$



One input, one output




- Let's add the bottom node now
- http://neuralnetworksanddeeplearning.com/chap4.html#universality_with_one_input_and_one_output

One input, one output



- Challenge: approximate this function




$$f(x) = 0.2 + 0.4x^2 + 0.3x \sin(15x) + 0.05 \cos(50x)$$

- Range and domain are $[0,1]$

One input, one output

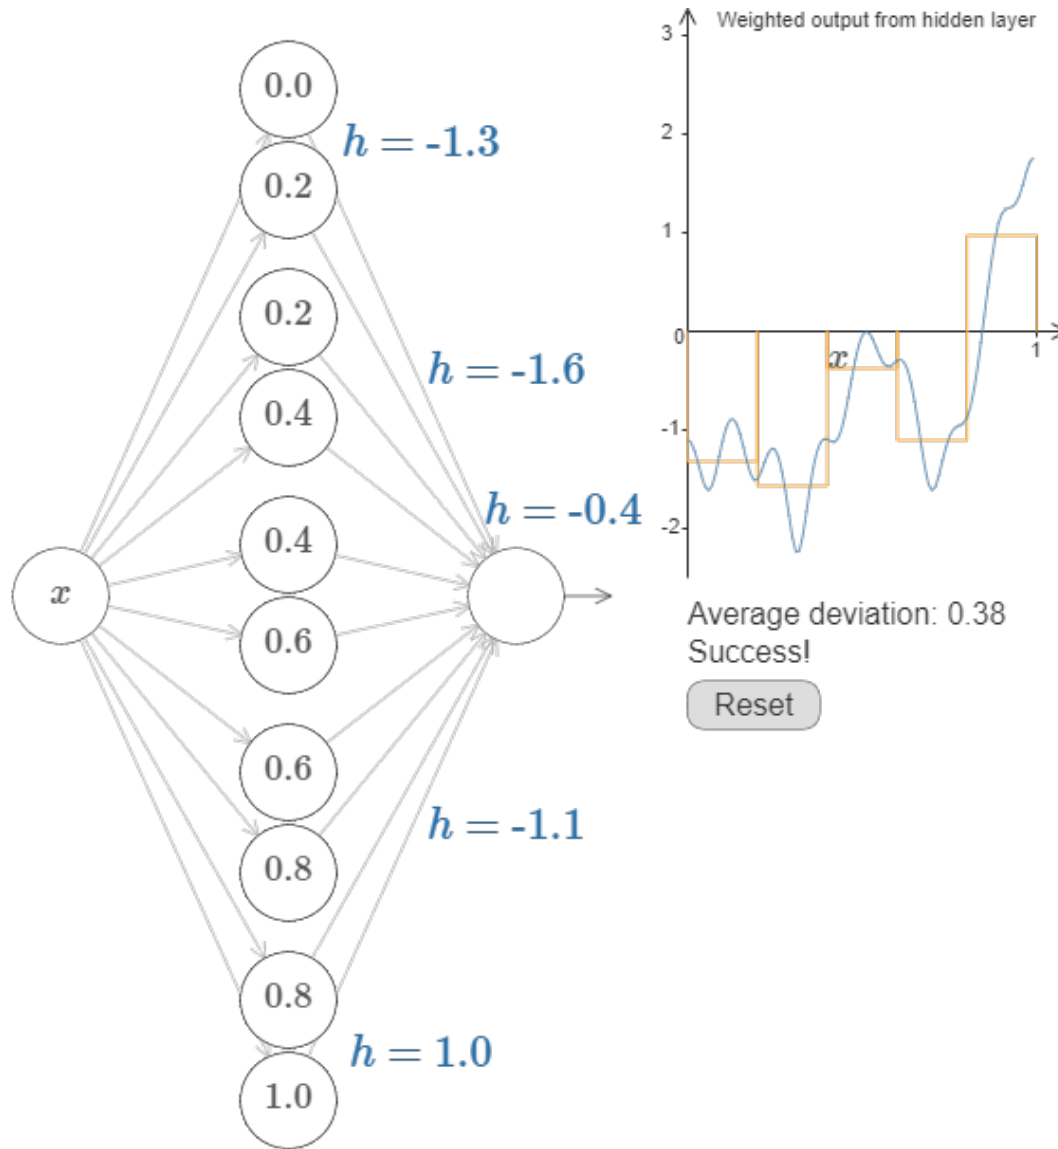


- We've been looking at the weighted combination from the hidden neurons $\sum_j w_j a_j$
- The actual output is $\sigma(\sum_j w_j a_j + b)$
- Take the inverse of the sigmoid function: $\sigma^{-1} \circ f(x)$



- http://neuralnetworksanddeeplearning.com/chap4.html#universality_with_one_input_and_one_output

One input, one output




One input, one output



- How do we convert back to standard parameterization?
 1. Set $w = 1000$ for first layer of weights
 2. Biases on hidden neurons are $b = -ws$
 3. Final layer of weights come from the $\pm h$ values
 4. Bias on the output neuron is 0

Many inputs

Two inputs




- http://neuralnetworksanddeeplearning.com/chap4.html#many_input_variables

Two inputs




- We've built something that looks a little like a tower function



Two inputs



- We can approximate arbitrary functions by adding towers of different heights in different locations



Two inputs



- With step functions, we've been implementing an if-then-else statement with neurons:

```
if input >= threshold:  
    output 1  
else:  
    output 0
```

- We can generalize this for multiple inputs:

```
if combined output from hidden neurons >= threshold:  
    output 1  
else:  
    output 0
```

- If we choose an appropriate threshold, we can squash the plateau down and leave only the tower

Two inputs




- Let's make a tower
- [http://neuralnetworksanddeeplearning.com/chap4.html#
many input variables](http://neuralnetworksanddeeplearning.com/chap4.html#many_input_variables)

Multiple outputs




- What about multiple outputs?
- A vector-valued function can be viewed as d real-valued functions
- We can simply construct a network approximating each component

Beyond sigmoid neurons





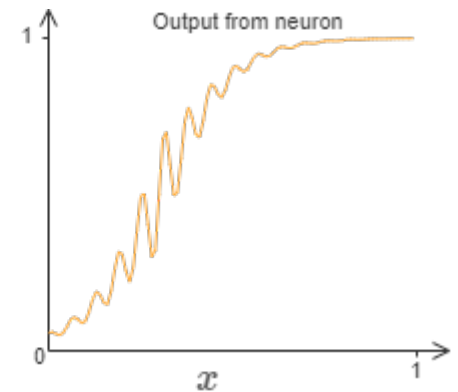



- What if we used this instead:



Beyond sigmoid neurons



- Increasing the weight gives an approximation of a step function
- Changing the bias changes the position of the step



Beyond sigmoid neurons

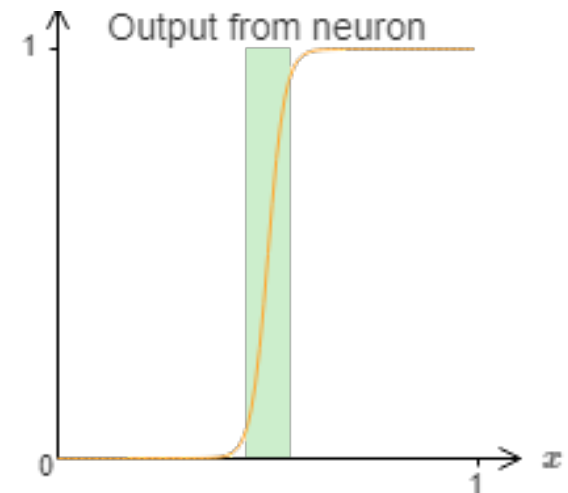


- What properties do we need for this approach?
 1. Need $s(z)$ to be well-defined as $z \rightarrow -\infty$ and $z \rightarrow \infty$
 - These are the values taken by the step function
 2. The limits must be different from each other
 - Otherwise we get a constant function
- These conditions are sufficient but not necessary for universality
 - The ReLU activation function also gives universality


Fixing the step functions



- We've been assuming our neurons produce exact step functions
- We actually only get an approximation
 - There's a narrow window of failure
 - We can increase the weights to make the window small
 - But is there a better way?




- Windows of failure have been exaggerated for demonstration





Fixing the step functions

- Let's approximate half the original function: $\sigma^{-1} \circ f(x)/2$
- Now approximate $\sigma^{-1} \circ f(x)/2$ shifted by half a bump
- Adding these together gives an overall approximation of $\sigma^{-1} \circ f(x)$
- The approximation is roughly a factor of 2 better in the windows of failure
- Could get further improvement by approximating $\sigma^{-1} \circ f(x)/M$ with M overlapping approximations



Wrap-up



- This explanation does not give a good prescription for designing neural networks!
 - Thus the result isn't directly useful for constructing networks
- However, universality answers the question of whether any particular function is computable with a neural network
- This changes the question to whether there is a good way to learn the function

Wrap-up



- If single layer network is universal, why use deep networks?
 - Note that our universality explanation required many hidden neurons
 - Earlier in the class, we argued that the hierarchical structure of deep networks is also helpful
- Summary:
 - Universality tells us neural networks can compute any function
 - Empirical evidence suggests deep networks are best adapted to learn those functions in practice

Further reading



- Nielsen book, chapter 4
- Roman Vershynin, High-Dimensional Probability, 2018